10 Scaling Laws

FOR PRODUCTION

RAG Systems

Most RAG systems fail not because of the LLM, but because of poor scaling decisions.

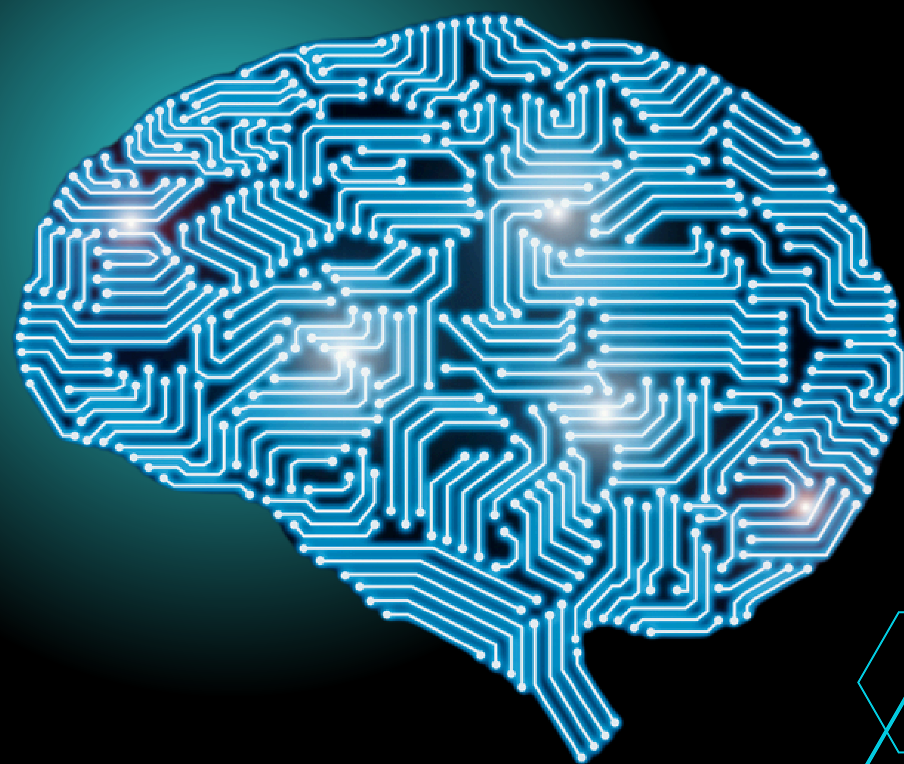

FOLLOW **Naresh Edagotti** to learn AI the easy way

LAW 1: DATA QUALITY DOMINATES QUANTITY

Principle

High-quality, well-structured data consistently outperforms large volumes of noisy information in RAG systems.

Production Techniques

- **Document Deduplication:** Use fuzzy matching (MinHash, SimHash) to identify near-duplicates
- **Content Quality Scoring:** Implement automated quality metrics (readability, factual density, source credibility)
- **Metadata Enrichment:** Add structured metadata (creation date, source, topic categories, confidence scores)
- **Domain-Specific Curation:** Establish subject matter expert review processes for critical domains

Implementation Checklist

- Automated duplicate detection pipeline
- Quality scoring thresholds and filtering
- Metadata schema standardization
- Regular content auditing workflows

LAW 2: CHUNKING STRATEGY DETERMINES RETRIEVAL SUCCESS

Principle

The granularity and method of document segmentation directly impacts both retrieval precision and context coherence.

Production Techniques

- **Semantic Chunking:** Split on natural boundaries (sentences, paragraphs, sections)
- **Sliding Window Approach:** Use overlapping chunks (typically 10-20% overlap)
- **Hierarchical Chunking:** Maintain document structure (headers, subsections)
- **Dynamic Chunk Sizing:** Adjust chunk size based on document type and domain

Optimal Chunk Sizes by Use Case

- **FAQ Systems:** 100-200 tokens
- **Technical Documentation:** 300-500 tokens
- **Research Papers:** 500-800 tokens
- **Legal Documents:** 200-400 tokens

LAW 3: EMBEDDING QUALITY IS NON-NEGOTIABLE

Principle

The semantic representation quality determines the upper bound of retrieval performance.

Production Techniques

- **Domain-Specific Fine-tuning:** Train embeddings on domain-specific corpora
- **Multi-Modal Embeddings:** Include text, metadata, and structural features
- **Embedding Versioning:** Track and manage embedding model updates
- **Periodic Re-Embedding:** Refresh embeddings as models improve

Recommended Models by Domain

- **General Purpose:** OpenAI text-embedding-3-large, Cohere embed-multilingual
- **Code/Technical:** CodeBERT, GraphCodeBERT
- **Legal:** Legal-BERT, CaseLaw-BERT
- **Medical:** BioBERT, ClinicalBERT

LAW 4: RETRIEVAL SCALES MORE EFFICIENTLY THAN MODEL SIZE

Principle
Expanding knowledge through improved retrieval is more cost-effective than using larger language models.

Production Techniques

- **Vector Database Optimization:**
  F AISS for CPU-intensive workloads
  Pinecone for managed cloud solutions
  Milvus for self-hosted scalability
- **Hybrid Retrieval:** Combine dense (semantic) and sparse (keyword) search
- **Approximate Nearest Neighbor (ANN):** Use HNSW, IVF, or LSH for speed
- **Distributed Indexing:** Shard large corpora across multiple nodes

Scaling Benchmarks

- **1M documents:** Single-node FAISS (sub-second queries)
- **10M documents:** Distributed Milvus cluster
- **100M+ documents:** Hybrid architecture with pre-filtering

LAW 5: RERANKING AMPLIFIES PRECISION

Principle

Initial retrieval casts a wide net; reranking refines results for relevance and quality.

Production Techniques

- **Cross-Encoder Reranking:** Use BERT-based models for precise relevance scoring
- **Multi-Stage Retrieval:** Fast retrieval (top 100) → precise reranking (top 10)
- **Learned Ranking:** Train custom rerankers on domain-specific query-document pairs
- **Ensemble Reranking:** Combine multiple ranking signals

Reranking Pipeline

- **First-Pass Retrieval:** Vector similarity (top 50-100 candidates)
- **Reranking:** Cross-encoder scoring (reduce to top 10-20)
- **Final Selection:** Context window optimization and deduplication

LAW 6: CONTEXT WINDOW ECONOMICS DRIVE ARCHITECTURE

Principle

Balancing context length with computational cost and latency requires strategic optimization.

Production Techniques

- **Dynamic Context Selection:** Vary context based on query complexity
- **Hierarchical Summarization:** Compress less relevant context
- **Multi-Query Retrieval:** Use query decomposition for complex questions
- **Context Caching:** Cache frequently accessed context segments

Context Optimization Strategies

- **Cost-Conscious:** 2K-4K token contexts with aggressive filtering
- **Quality-Focused:** 8K-16K token contexts with light compression
- **Hybrid Approach:** Dynamic adjustment based on query confidence

LAW 7: FRESHNESS BEATS STATIC KNOWLEDGE

Principle

Dynamic, up-to-date knowledge bases outperform larger but stale datasets.

Production Techniques

- **Incremental Indexing:** Real-time document ingestion and embedding
- **Time-to-Live (TTL) Policies:** Automatic expiration of outdated content
- **Version Control:** Track document changes and maintain history
- **Scheduled Refresh:** Periodic re-crawling and re-embedding

Freshness Implementation

- **Real-time:** Stream processing for critical updates (news, prices)
- **Near Real-time:** Batch processing every 5-15 minutes
- **Periodic:** Daily/weekly updates for stable content

LAW 8: LATENCY-ACCURACY TRADEOFFS ARE INEVITABLE

Principle

More sophisticated retrieval improves accuracy but increases latency—optimize for your use case.

Production Techniques

- **Multi-Tier Caching:**
  L1: In-memory query results (milliseconds)
  L2: Precomputed embeddings (sub-second)
  L3: Full retrieval pipeline (1-3 seconds)

- **Hybrid Retrieval:** Combine dense (semantic) and sparse (keyword) search
- **Approximate Nearest Neighbor (ANN):** Use HNSW, IVF, or LSH for speed
- **Distributed Indexing:** Shard large corpora across multiple nodes

FLatency Targets by Application

- **Chatbots:** <2 seconds end-to-end
- **Search Systems:** <500ms retrieval + <1s generation
- **Analytics:** <10s for complex queries

LAW 9: EVALUATION MUST SCALE WITH SYSTEM COMPLEXITY

Principle

Traditional NLP metrics (BLEU, ROUGE) are insufficient for production RAG systems.

Production Techniques

- **Faithfulness Metrics:** Measure grounding in retrieved documents
- **Hallucination Detection:** Automated fact-checking against sources
- **Human-in-the-Loop:** Regular expert evaluation of edge cases
- **A/B Testing:** Compare system versions with real user interactions

Evaluation Framework

1. **Retrieval Quality:** Precision@K, Recall@K, MRR
2. **Generation Quality:** Faithfulness, completeness, coherence
3. **System Performance:** Latency, throughput, error rates
4. **User Satisfaction:** Click-through rates, task completion

LAW 10: SAFETY AND GUARDRAILS ARE NON-OPTIONAL

Principle

Retrieved content can introduce bias, misinformation, or inappropriate material into generated responses.

Production Techniques

- **Content Filtering:** Pre-ingestion toxicity and bias detection
- **Source Verification:** Credibility scoring and domain whitelisting
- **Prompt Injection Defense:** Sanitize retrieved content before LLM input
- **Output Validation:** Post-generation fact-checking and safety filters

Safety Pipeline

1. **Ingestion:** Content filtering and source validation
2. **Retrieval:** Query sanitization and result filtering
3. **Generation:** Prompt protection and output validation
4. **Monitoring:** Continuous safety metric tracking

FOLLOW  **Naresh Edagotti** to learn AI the easy way

KEY TAKEAWAYS

✅ Quality over quantity in data curation
✅ Strategic chunking based on content type and use case
✅ Multi-stage retrieval with reranking for precision
✅ Dynamic optimization of context and caching
✅ Comprehensive evaluation beyond traditional metrics
✅ Proactive safety measures throughout the pipeline

Remember:
Scaling RAG isn't about adding more documents—it's about building smarter, more efficient systems that deliver better results with optimal resource utilization.

LIKE THIS CONTENT?

FOLLOW FOR MORE!

NARESH EDAGOTTI

LIKE REPOST SAVE