


Ace Gen-AI Interviews: 25 Q&A

Essential Interview Questions
and Expert Answers for Your
Generative AI Career



2. What is the Attention Mechanism in Transformers?


Attention mechanisms allow the model to focus on relevant parts of the input. Key types include:




These mechanisms directly influence generation quality and speed by enabling efficient handling of dependencies.




Transformers vs. RNNs: Why Transformers Win




Transformer-Based Models: Architectures & Variants



What is a Large Language Model (LLM)?



LLM Capabilities & Challenges: A Trade-off Analysis



Hallucination in LLMs


When an LLM produces plausible yet false statements.
Critical for evaluation and safeguards.

Definition



Plausible but incorrect outputs, often due to extrapolation.

Primary Causes



Importance & Action



Evaluating Outputs



Designing Safeguards

Crucial for reliable AI systems, involving fact-checking, RAG, and human oversight.

Hallucinations in LLMs: Creativity vs. Factuality

Fueling Creativity (Non-Factual Domains)



Storytelling & Brainstorming




Idea Generation




Exploring Narrative Possibilities

Mitigation Strategies (Factual Domains)




Prompt Engineering: Crafting Inputs for Desired Outputs

Prompt engineering is the art of designing input text to guide Large Language Models (LLMs) towards specific, high-quality responses. It involves strategic phrasing and structuring of instructions.




Common Prompting Techniques & When to Use Them



Selecting the right pattern aligns model behavior with task complexity and reliability needs.

What is Retrieval-Augmented Generation (RAG)?



RAG couples an LLM with a retriever that fetches relevant documents at inference time. This grounds responses in fresh, authoritative context, cutting hallucination and eliminating costly retraining when knowledge changes.



Choosing the Right Approach: Prompting vs RAG vs Fine-Tuning

Strategic Selection for Generative AI Implementation




Prompt Engineering

- Quick Prototyping
- Small Context Tasks
- Rapid Iteration
- Low Resource Needs




Retrieval-Augmented Generation (RAG)

- Up-to-date Information
- Proprietary Data Integration
- Grounded Retrieval
- Reduces Hallucination




Fine-Tuning

- Deep Domain Adaptation
- Style Mimicry
- Specialized Behavior
- High Customization



Key Takeaway: Align the approach with specific task requirements, data availability, and desired model behavior for optimal results.


LangChain: Modular Framework for LLM Applications




LangChain abstracts and modularizes complex LLM interactions, enabling rapid development of advanced, connected, and autonomous AI applications.

14. What is LlamaIndex?


Connects LLMs with Private Data for **Natural Language Querying**.



Fine-tuning in LLMs: Transforming General Models into Bespoke Experts




Ace Gen-AI Interviews: Fine-Tuning Techniques & Optimization



17. What is LoRA in Fine-Tuning & QLoRA?

LoRA: Low-Rank Adaptation




Injects trainable low-rank matrices into transformer layers.

Enables task adaptation with <1% extra weights.

Minimal resource use for fine-tuning.

QLoRA: Quantized LoRA



Quantizes base weights to 4 bits.

Further cuts memory footprint.


Combines 4-bit quantization with LoRA for highly efficient tuning.

Benefits: Democratizing Customization


- Delivers near-full fine-tune quality.
- Runs efficiently on consumer GPUs.
- Makes model customization accessible to a wider audience.

Small Language Models (SLMs): Efficient & On-Device

Millions to a few billion parameters for optimized performance.



Diffusion Models: Reversing Noise to Generate




Key Advantages over GANs:

- Iterative refinement for high fidelity.
- Better diversity and stability.




Noise Schedules in Diffusion Models

The noise schedule dictates variance at each diffusion timestep, controlling the data corruption process.



Strategic noise schedule design is crucial for efficient and high-fidelity generative processes.

Multimodal LLMs



Ingest and produce across multiple modalities within a single, unified architecture.

Examples: GPT-4V, DALL·E, Flamingo

Applications: Visual Question Answering (VQA), Image Captioning, Cross-modal Retrieval, Expanding Generative AI beyond text.

Common Challenges in Deploying Generative AI Systems

1. Hallucination & Fact-Checking



Confidently generated but factually incorrect responses. Requires rigorous verification and grounding strategies.

2. Prompt Sensitivity



Outputs highly dependent on precise input phrasing. Necessitates extensive prompt engineering and iterative testing.

3. Data Privacy & Compliance



Ensuring user data protection and adhering to regulatory standards (e.g., GDPR, HIPAA) across the AI lifecycle.

4. High Computational Costs



Training and inference require substantial computational resources, leading to significant operational expenses.

5. Evaluation Complexity & Benchmarks



Lack of standardized benchmarks makes objective performance and reliability assessment difficult.

Anticipating these challenges early informs architecture choices, governance policies, and budget planning for successful AI deployment.

Mitigating LLM Hallucinations: A Critical Challenge

Unchecked hallucination erodes user trust. Safeguarding brand reputation and regulatory adherence is paramount.



HALLUCINATION RISK

Confidence Thresholds: Establish clear limits to prevent presenting unsure information.

Fallback Responses: Design pre-defined responses for uncertain scenarios.



MITIGATION STRATEGIES



RAG (Retrieval-Augmented Generation): Ground responses in verified knowledge bases.



Uncertainty Scoring: Implement mechanisms to quantify response confidence.



Real-Time Verification APIs: Utilize external tools for fact-checking.



Human-in-the-Loop Review: Involve human experts for complex cases.

By implementing these strategies, organizations can minimize risk and build reliable Gen-AI systems.

Data Privacy and Compliance Challenges



Key Regulations & Mandates

- **GDPR & HIPAA:** Mandate strict data handling.
- **Data Minimization:** Collect only necessary data.
- **User Consent:** Obtain explicit permission.
- **Right-to-be-Forgotten:** Allow data deletion.



Key Techniques for Compliance

- **On-Prem Deployment:** Keep data locally secure.
- **Differential Privacy:** Add noise to protect individuals.
- **Synthetic Data Generation:** Use fake data for training.
- **Strict Access Controls:** Limit data access.




The Privacy-by-Design Advantage


- **Proactive Embedding:** Integrate privacy early.
- **Avoids Retrofits:** Prevents costly later changes.
- **Mitigates Legal Exposure:** Reduces compliance risks.
- **Builds Trust:** Enhances user confidence.

Embedding privacy-by-design principles avoids costly retrofits and significant legal exposure.

Gen-AI Cost Drivers & Optimization Levers




Generative AI Deployment: Challenges & Evaluation Frameworks




Strategic Mastery & Durable Competitive Advantage in Generative AI

Balancing Technical Depth, Deployment Constraints, and Core Requirements for Career Success


TECHNICAL MASTERY & TRADE-OFFS



CORE REQUIREMENTS & MITIGATION



STRATEGIC ADVANTAGE & ADAPTATION



Integrating these pillars secures a durable career foundation in the evolving landscape of Generative AI applications.